



Spring Batch Performance Tuning

By Chris Schaefer & Gunnar Hillert

Agenda

- Spring Batch
- Spring Integration
- Spring Batch Integration
- Scaling Spring Batch
- Spring XD

Spring Batch

<http://projects.spring.io/spring-batch/>



Batch processing ... is defined as the processing of data without interaction or interruption.

Michael T. Minella, Pro Spring Batch

Batch Jobs

- Generally long-running
- Non-interactive
 - Often include logic for handling errors and restartability options
- Process large volumes of data
 - More than what may fit in memory or a single transaction

Batch and offline processing

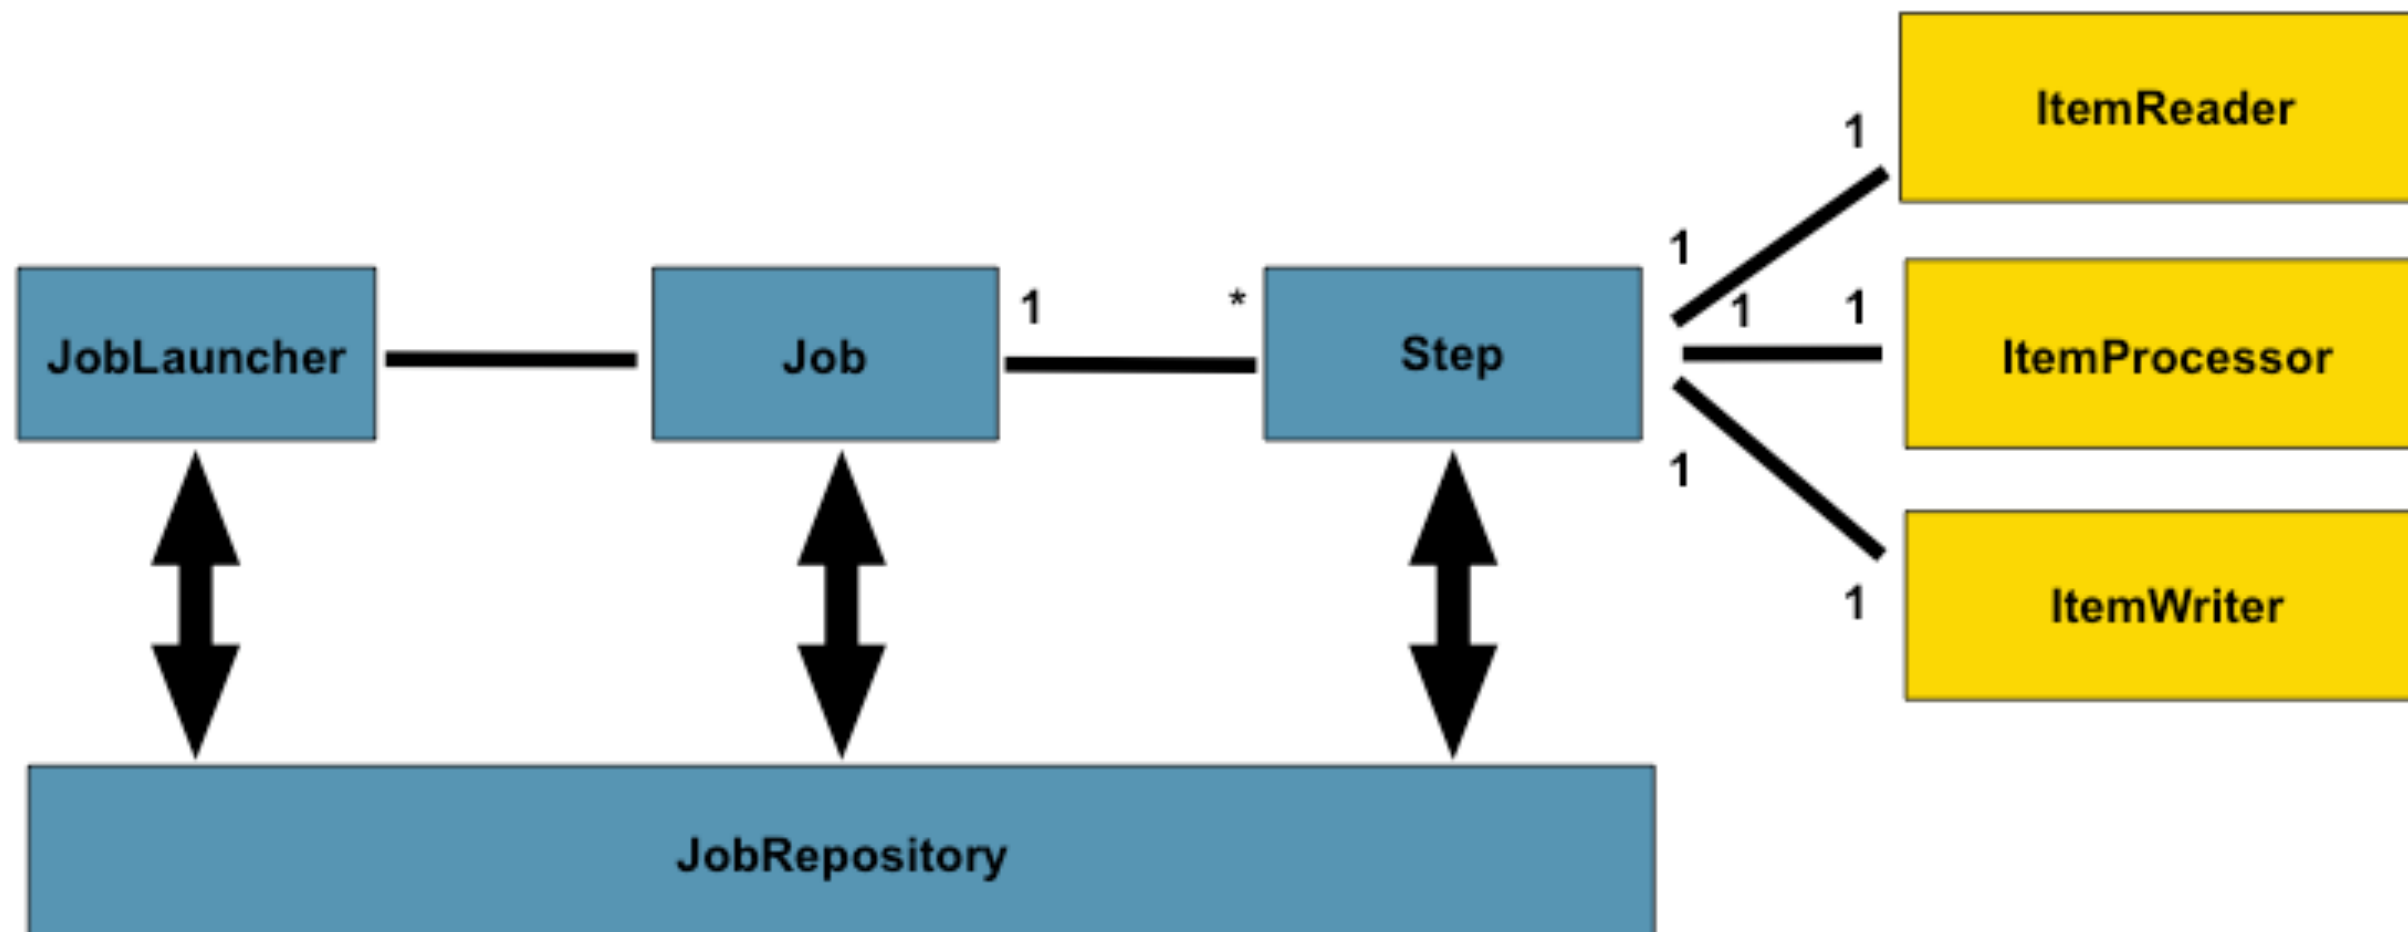
- Close of business processing
 - Order processing, Business reporting, Account reconciliation, Payroll
- Import / export handling
 - a.k.a. ETL jobs (Extract-Transform-Load)
 - Data warehouse synchronization
- Large-scale output jobs
 - Loyalty program emails, Bank statements
- Hadoop job orchestration

Features

- Transaction management
- Chunk based processing
- Schema and Java Config support
 - Annotations for callback type scenarios such as Listeners
- Start/Restart/Skip capabilities
- Based on the Spring framework
- JSR 352: Batch Applications for the Java Platform

Concepts

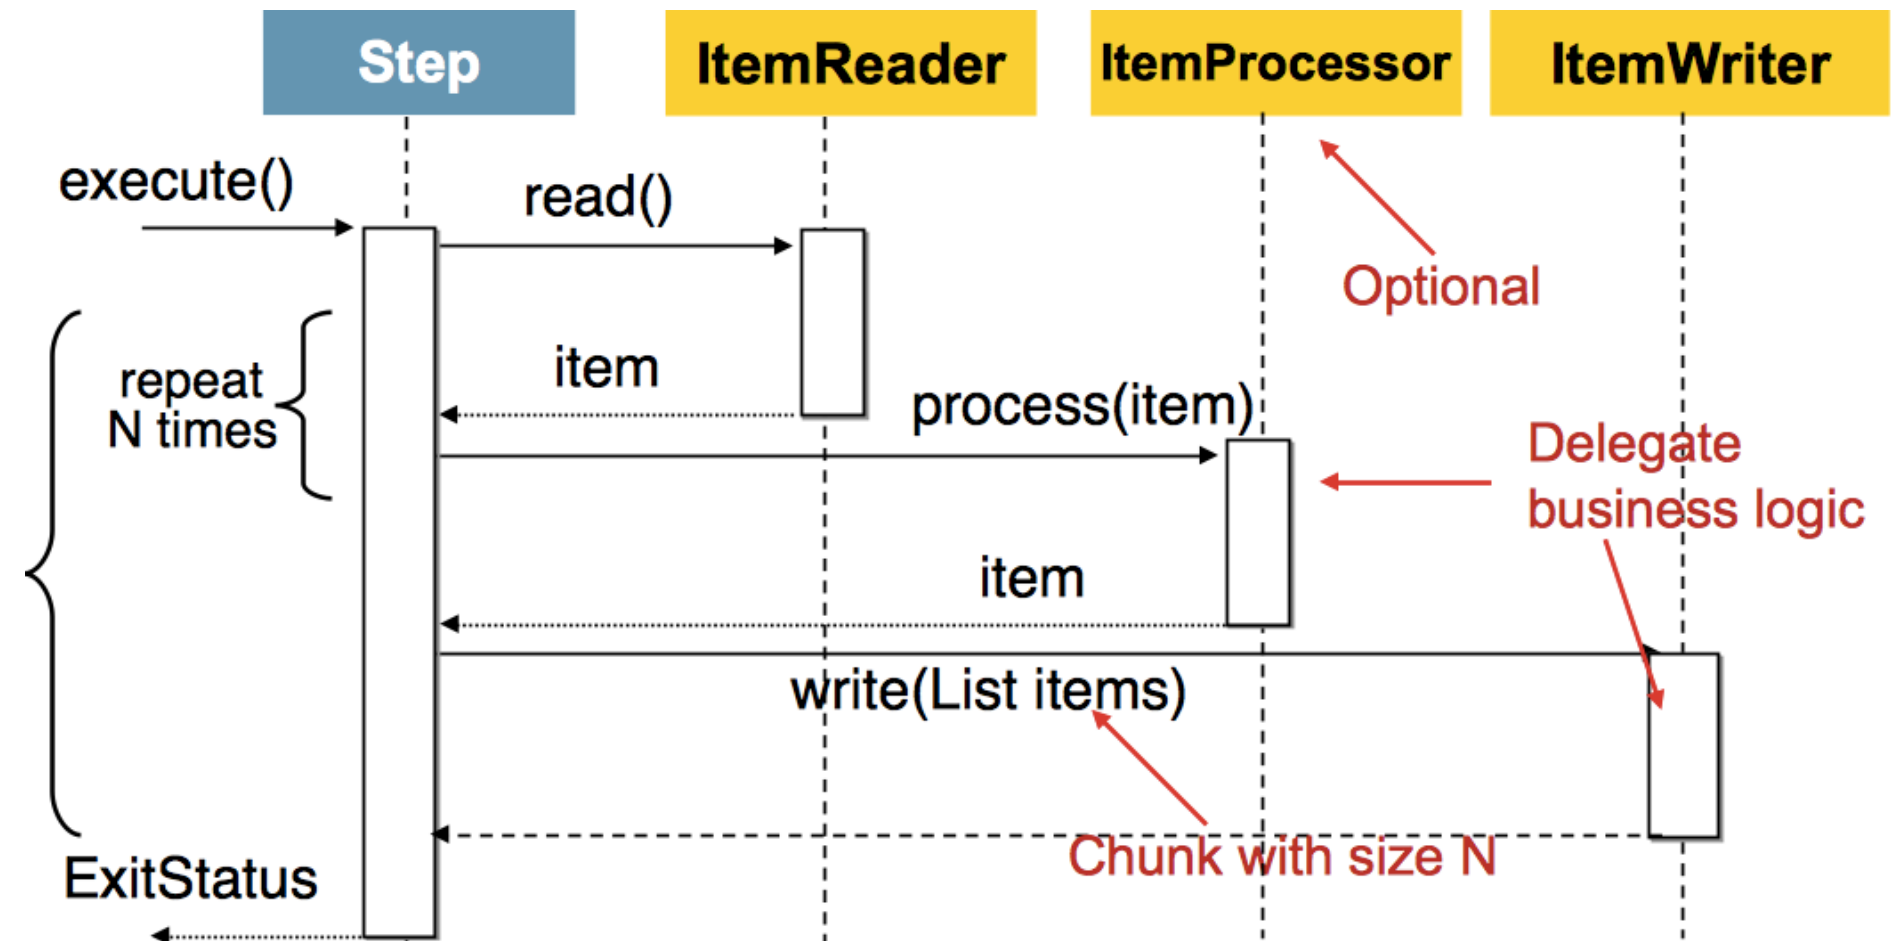
- Job
- Step
- Chunk
- Item



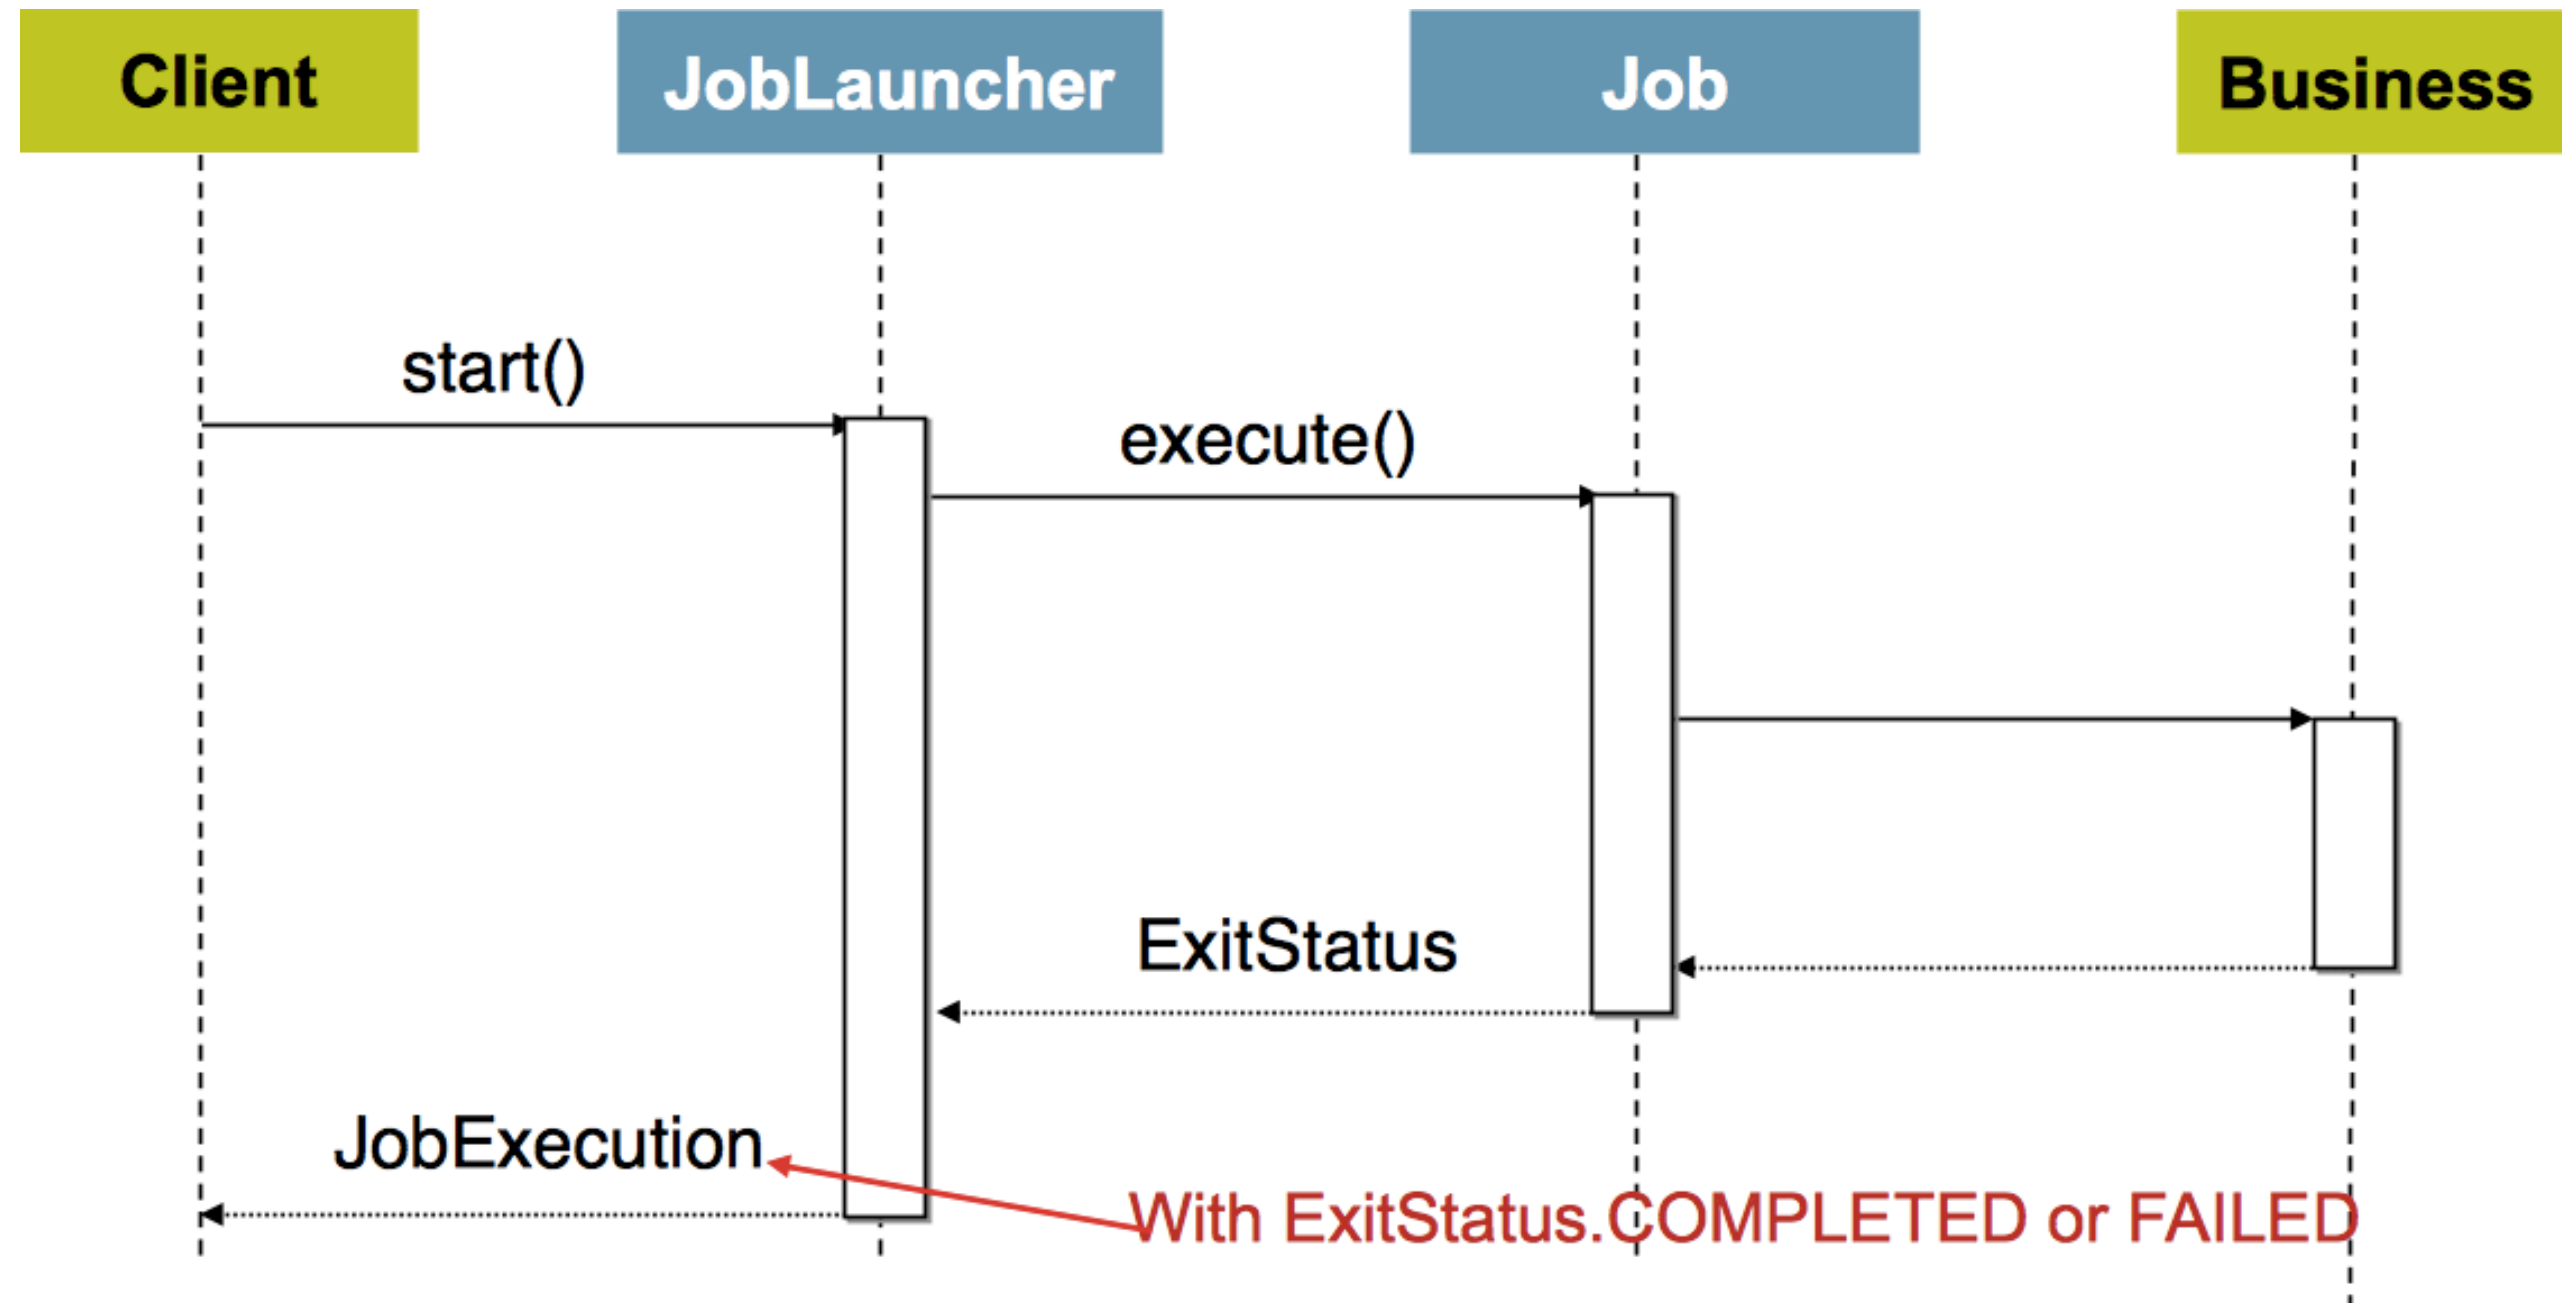
Repeat | Retry | Skip | Restart

Chunk-Oriented Processing

- Read data, optionally process and write out the “chunk” within a transaction boundary.



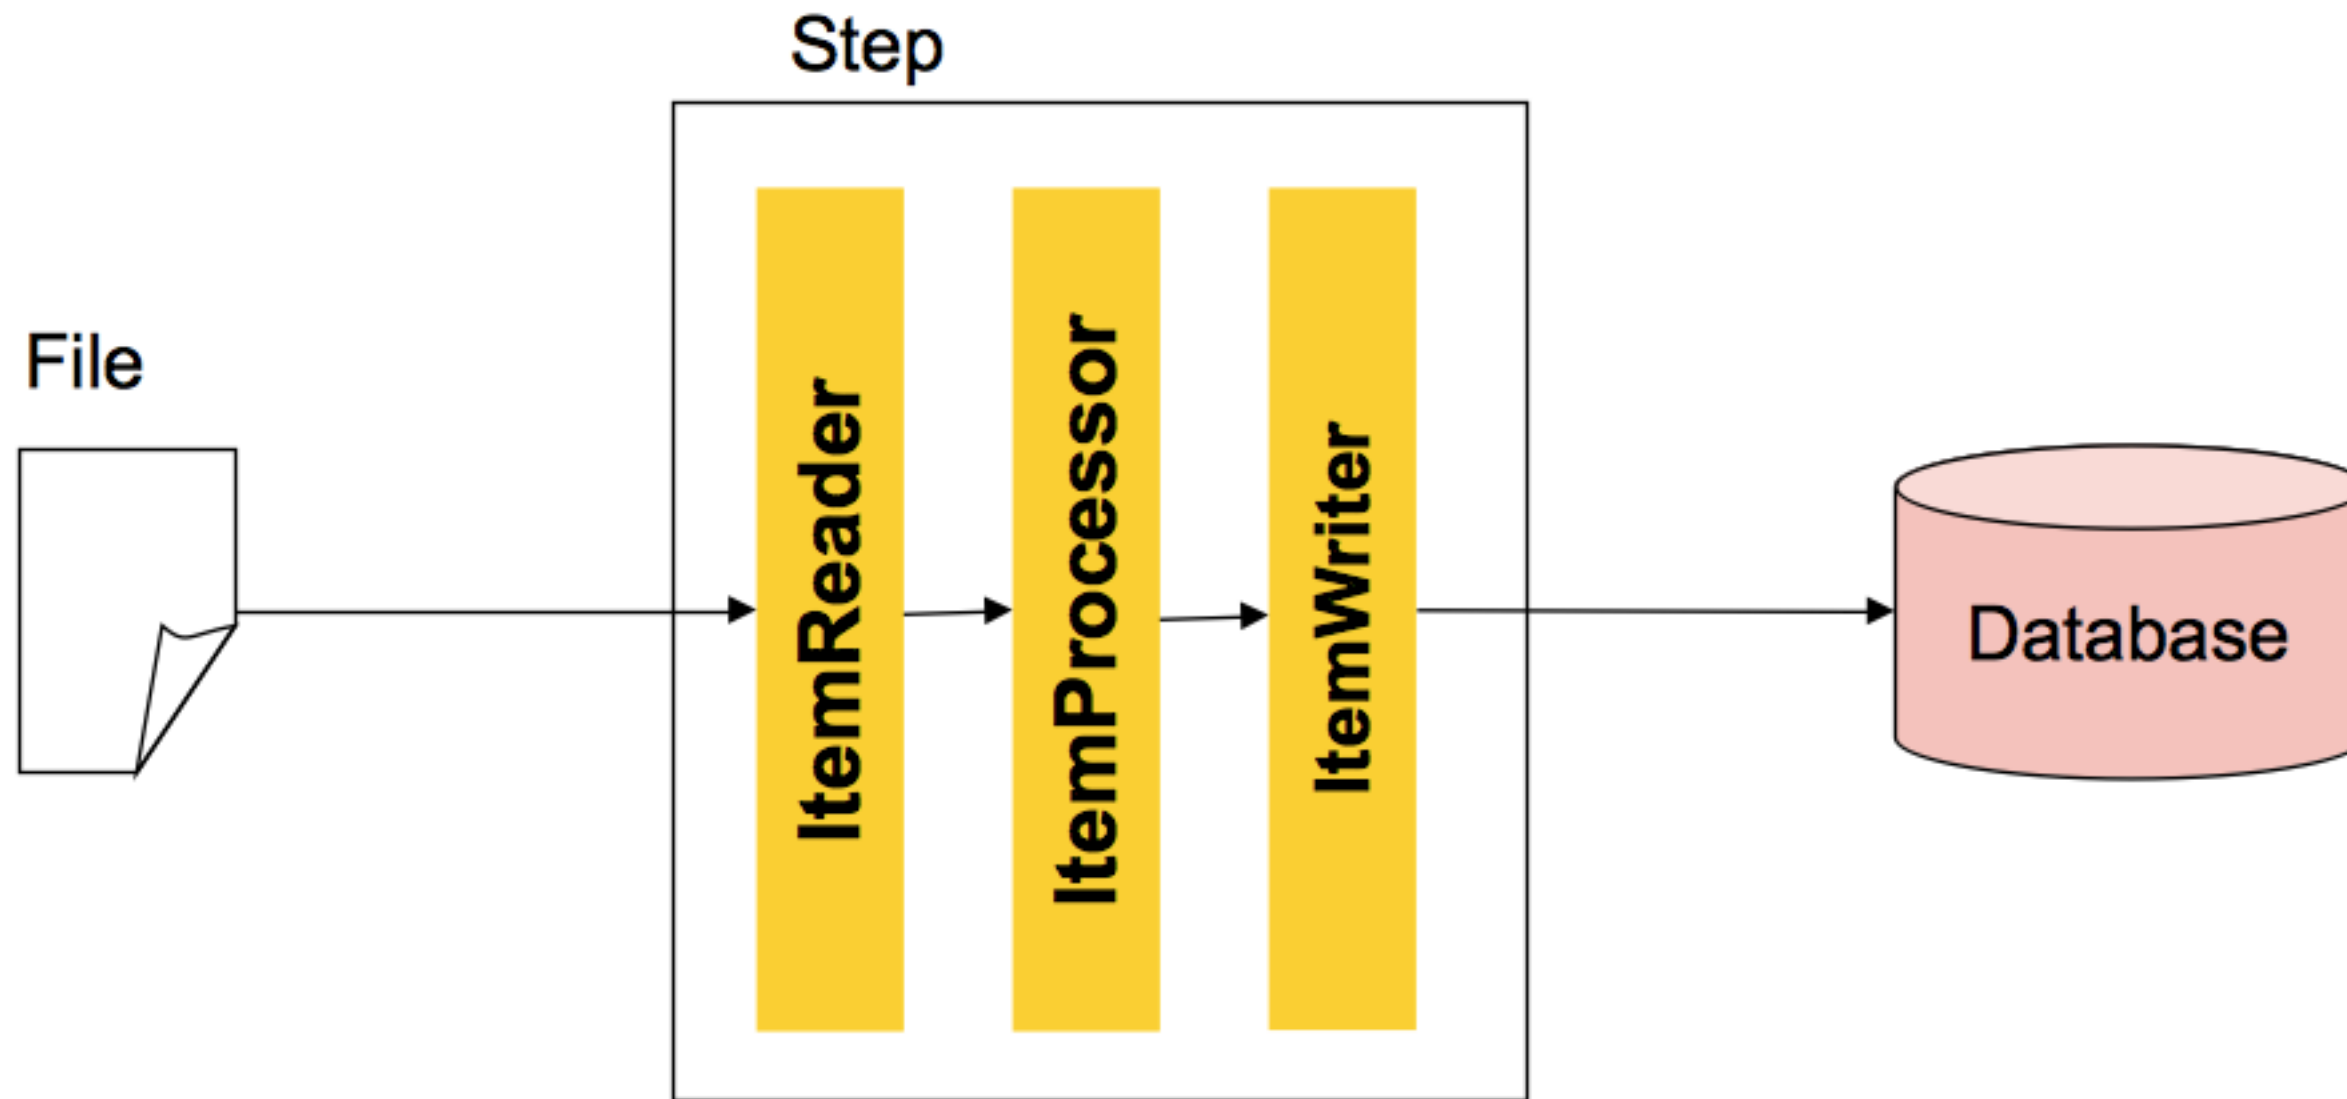
JobLauncher



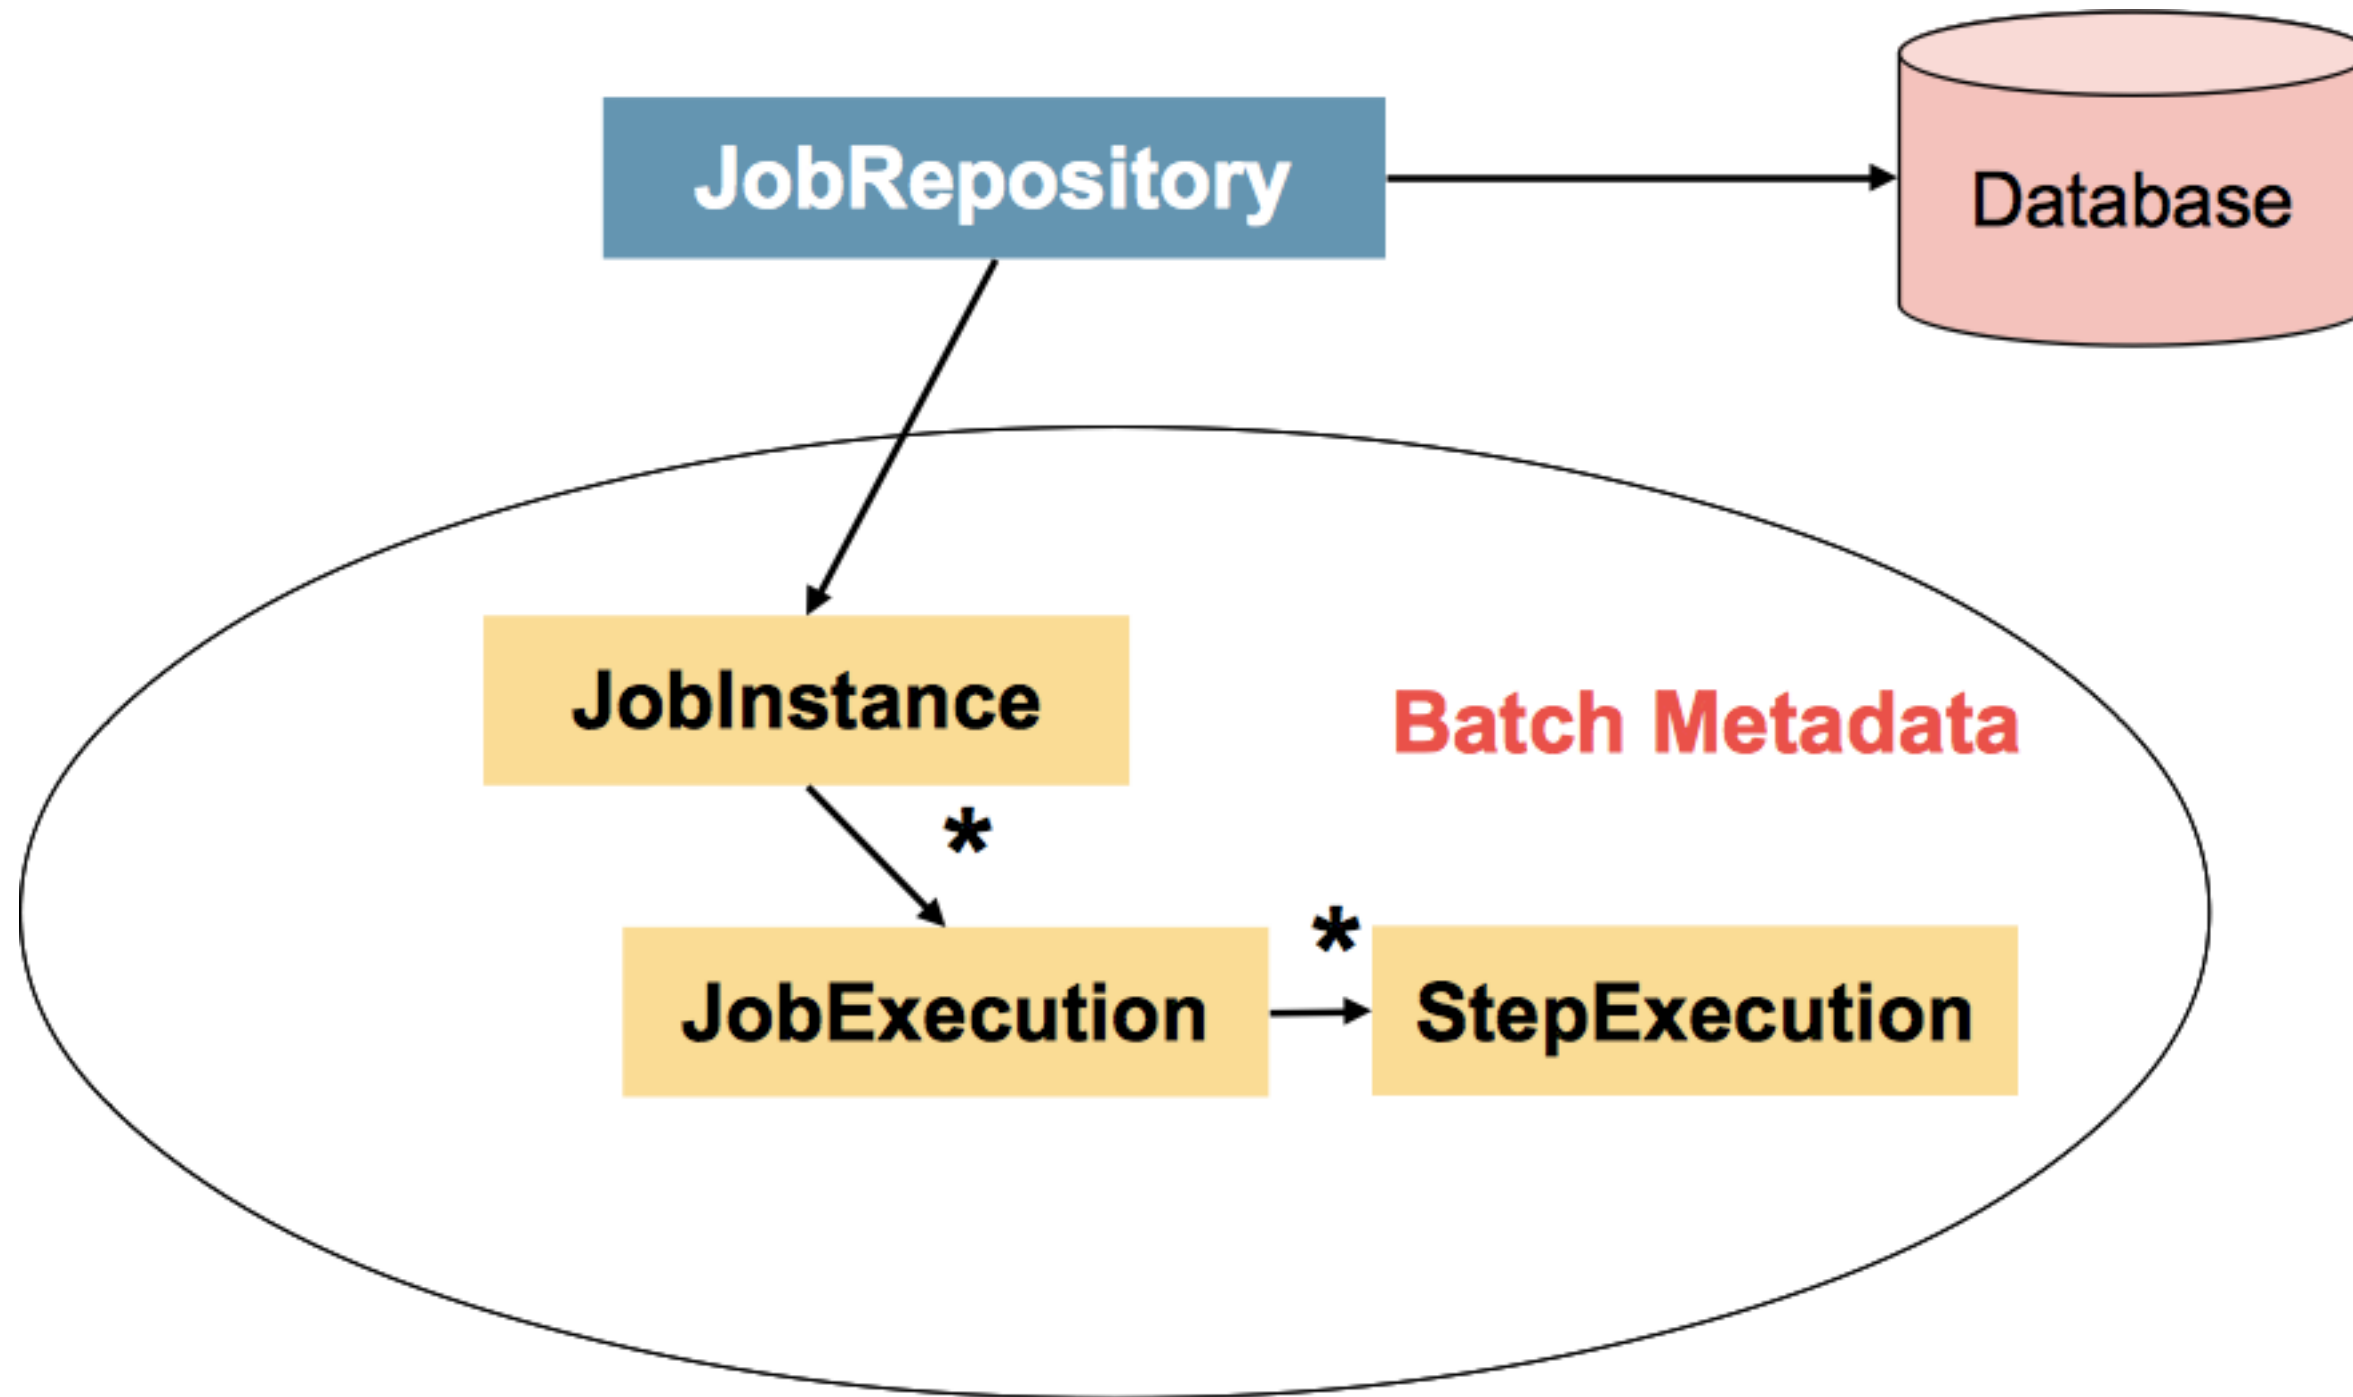
ItemReaders and ItemWriters

- Flat File
- XML (StAX)
- Multi-File Input
- Database
- JDBC, JPA/Hibernate, Stored Procedures, Spring Data
- JMS
- AMQP
- Email
- Implement your own...

Simple File Load Job



Job Repository



Spring Integration

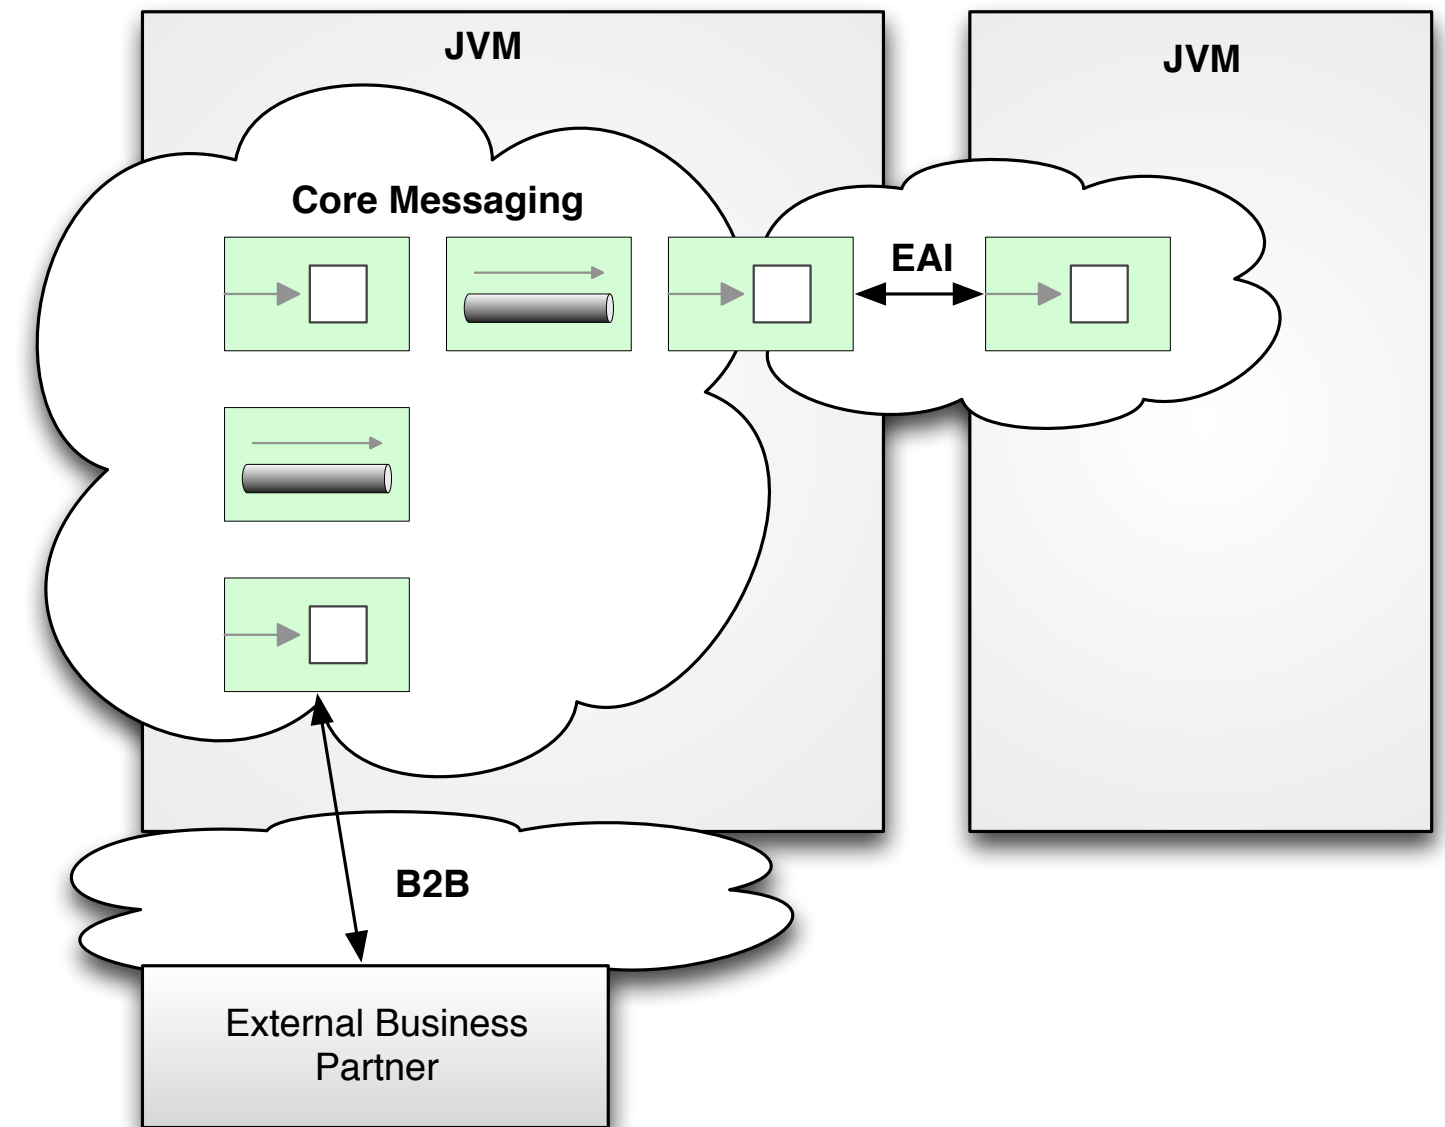
<http://projects.spring.io/spring-integration/>

Integration Styles

- File Transfer
- Shared Database
- Remoting
- Messaging

Integration Styles

- Business to Business Integration (B2B)
- Inter Application Integration (EAI)
- Intra Application Integration




Common Patterns




Enterprise Integration Patterns

- By Gregor Hohpe & Bobby Woolf
- Published 2003
- Collection of well-known patterns
- Icon library provided



<http://www.eaipatterns.com/eaipatterns.html>





Spring Integration provides an extension of the Spring programming model to support the well-known enterprise integration patterns.

Spring Integration Website

Adapters



AMQP/RabbitMQ
AWS
File/Resource
FTP/FTPS/SFTP
GemFire
HTTP (REST)
JDBC
JMS
JMX
JPA

MongoDB
POP3/IMAP/SMTP
Print
Redis
RMI
RSS/Atom
SMB
Splunk
Spring Application
Events


Stored Procedures
TCP/UDP
Twitter
Web Services
XMPP
XPath
XQuery

Custom Adapters

Samples

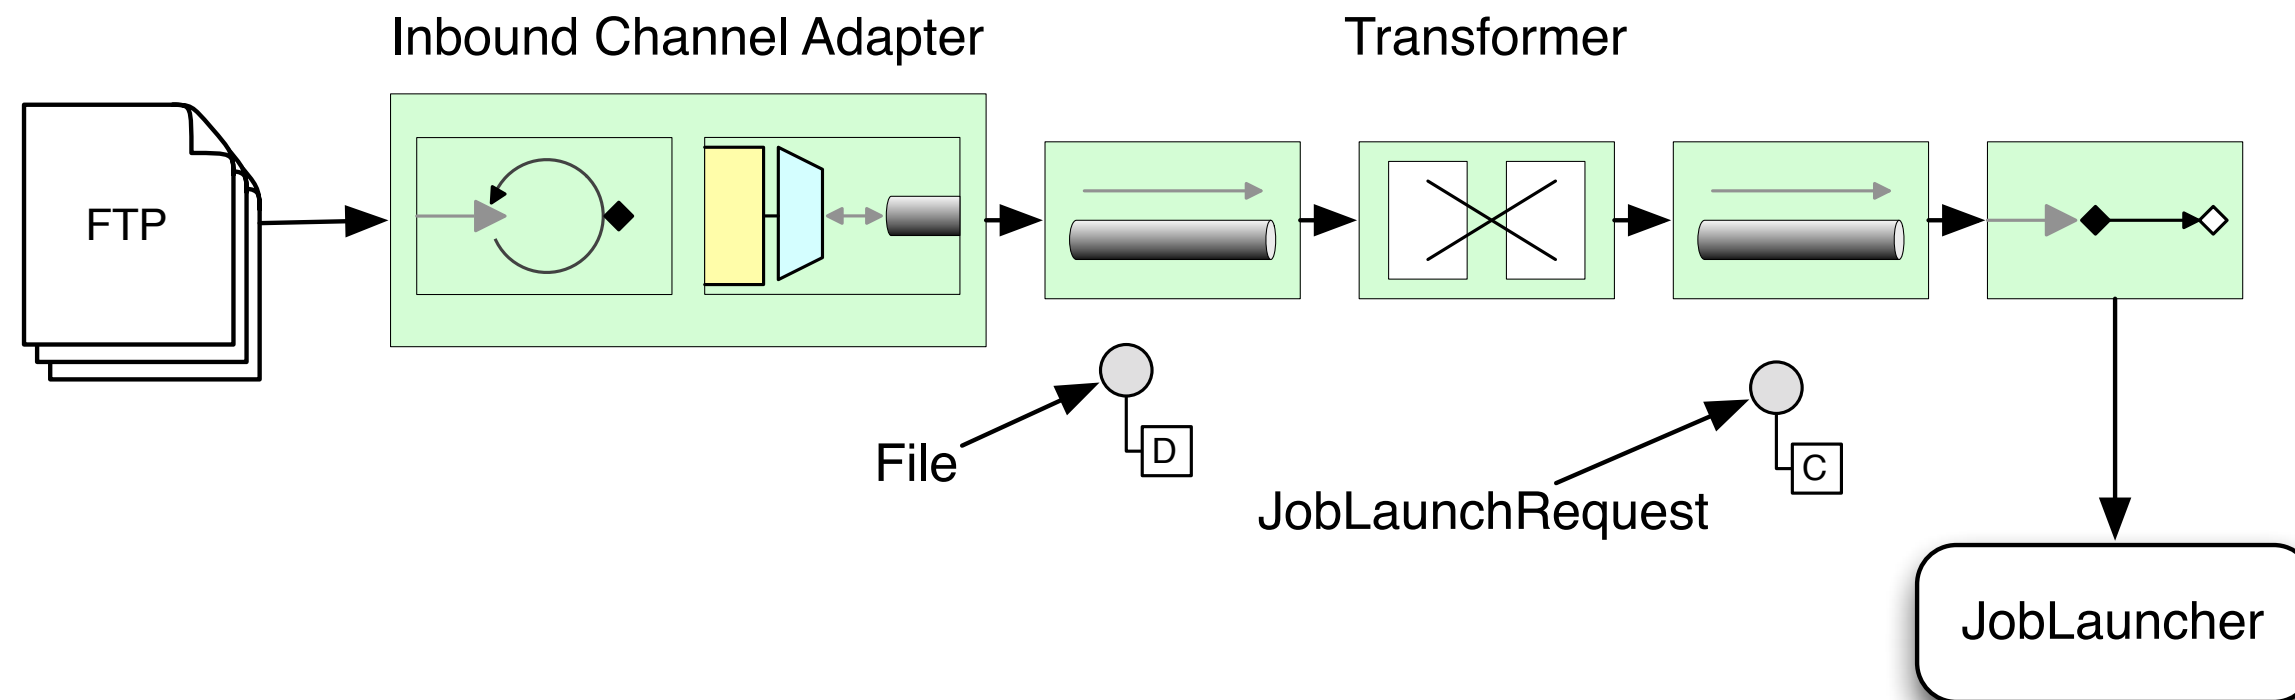
- <https://github.com/spring-projects/spring-integration-samples>
- Contains 50 Samples and Applications
- Several Categories:
 - Basic
 - Intermediate
 - Advanced
 - Applications

Spring Batch Integration



Launching batch jobs through messages

- Event-Driven execution of the JobLauncher
- Spring Integration retrieves the data (e.g. file system, FTP, ...)
- Easy to support separate input sources simultaneously



JobLaunchRequest



```
public class FileMessageToJobRequest {  
    private Job job;  
    private String fileParameterName;  
    ...  
    @Transformer  
    public JobLaunchRequest toRequest(Message<File> message) {  
        JobParametersBuilder jobParametersBuilder = new JobParametersBuilder();  
        jobParametersBuilder.addString(fileParameterName,  
                                       message.getPayload().getAbsolutePath());  
        return new JobLaunchRequest(job, jobParametersBuilder.toJobParameters());  
    }  
}
```

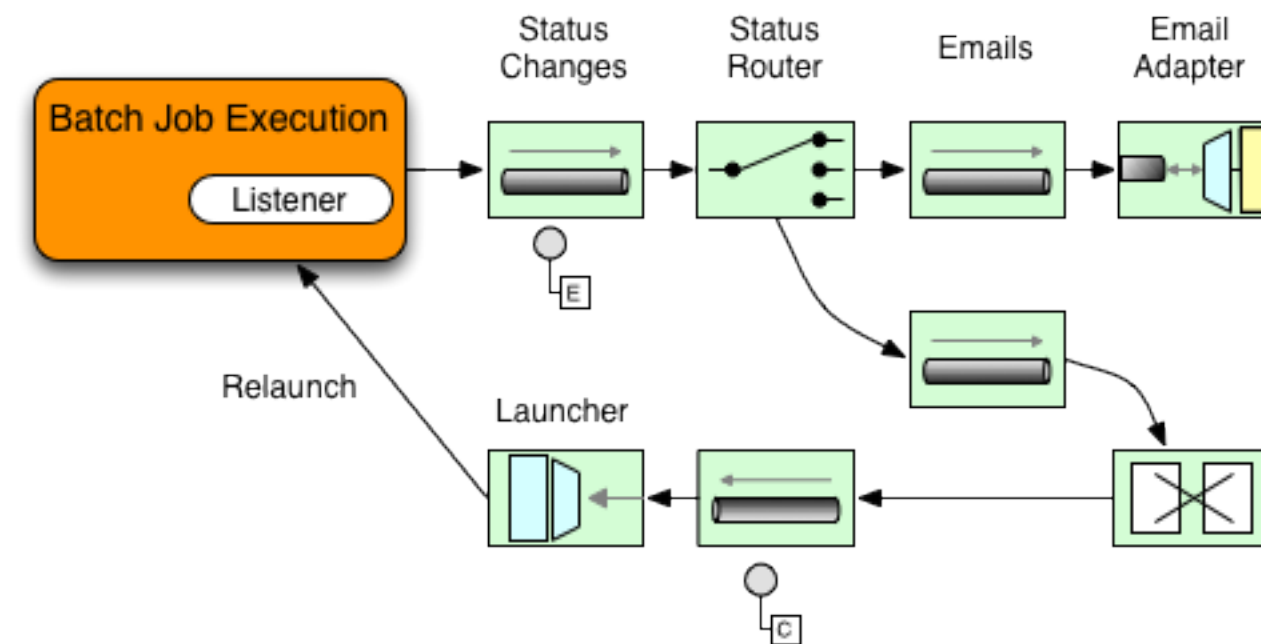

JobLaunchRequest



```
<batch-int:job-launching-gateway    request-channel="requestChannel"  
    reply-channel="replyChannel"  
    job-launcher="jobLauncher" />
```

Get feedback with informational messages

- Spring Batch provides support for listeners:
 - StepExecutionListener
 - ChunkListener
 - JobExecutionListener



Get feedback with informational messages

```
<batch:job id="importPayments">  
  ...  
  <batch:listeners>  
    <batch:listener ref="notificationExecutionsListener"/>  
  </batch:listeners>  
</batch:job>  
  
<int:gateway id="notificationExecutionsListener"  
  service-interface="o.s.batch.core.JobExecutionListener"  
  default-request-channel="jobExecutions"/>
```

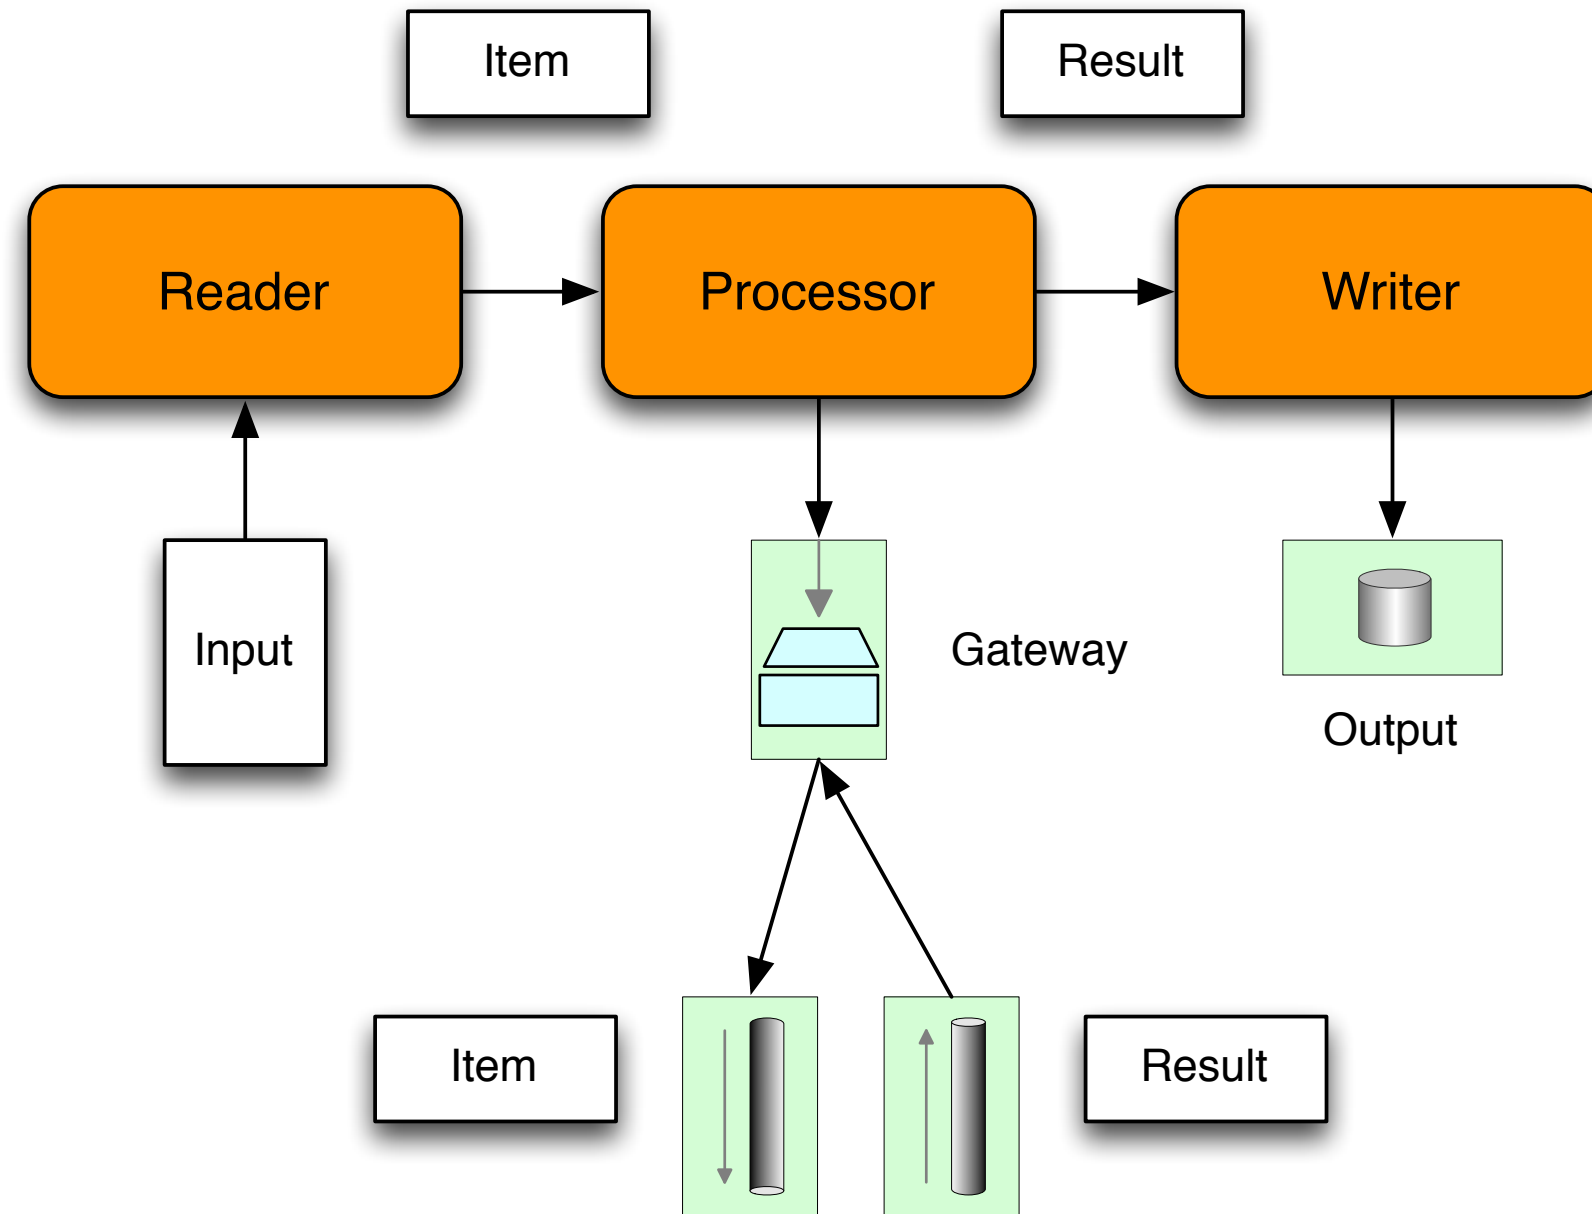
Launching and information messages demo in next section

Scaling Spring Batch

Scaling and externalizing batch process execution

- Utilization of Spring Integration for multi process communication
- Distribute complex processing
 - Single process
 - Multi-threaded steps
 - Parallel steps
 - Local partitioning
 - Multi process
 - Remote chunking
 - Remote partitioning
- Asynchronous Item processing support
 - AsyncItemProcessor
 - AsyncItemWriter

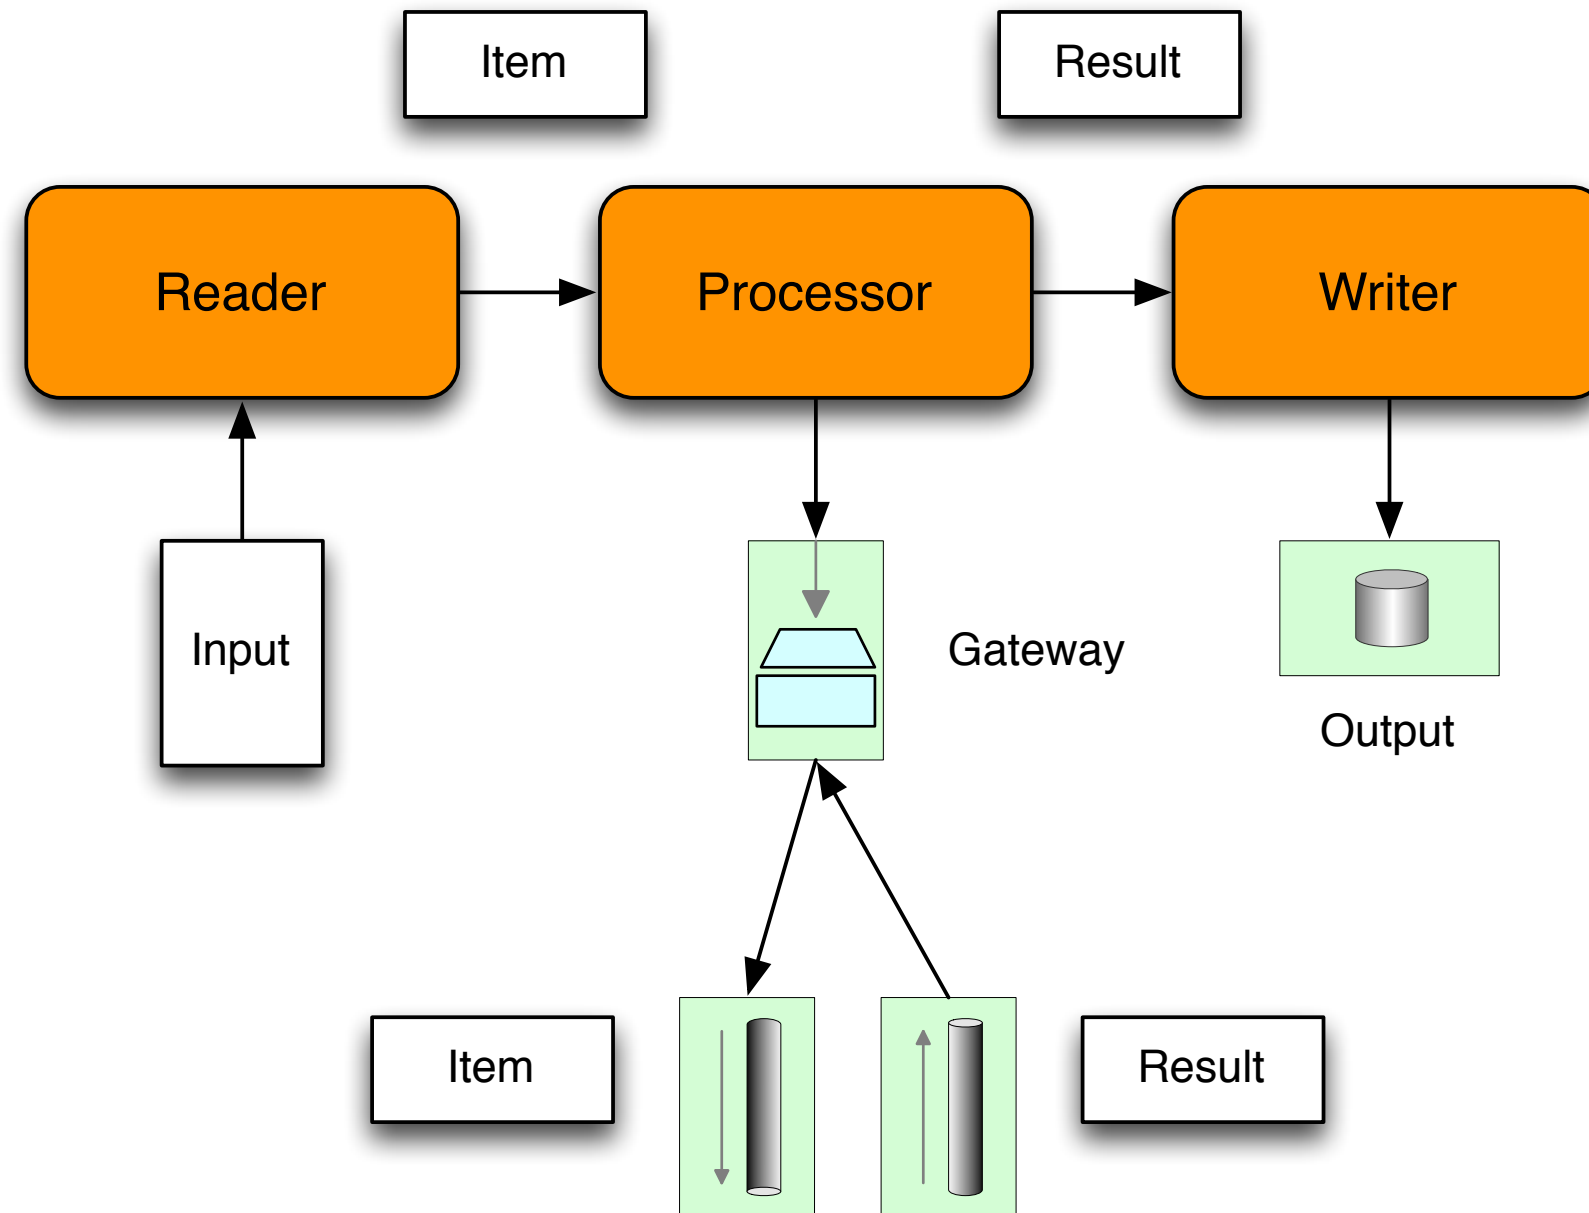
Single Thread



Single Thread - Demo

Multi-threaded

- Simply add a TaskExecutor to your Tasklet configuration



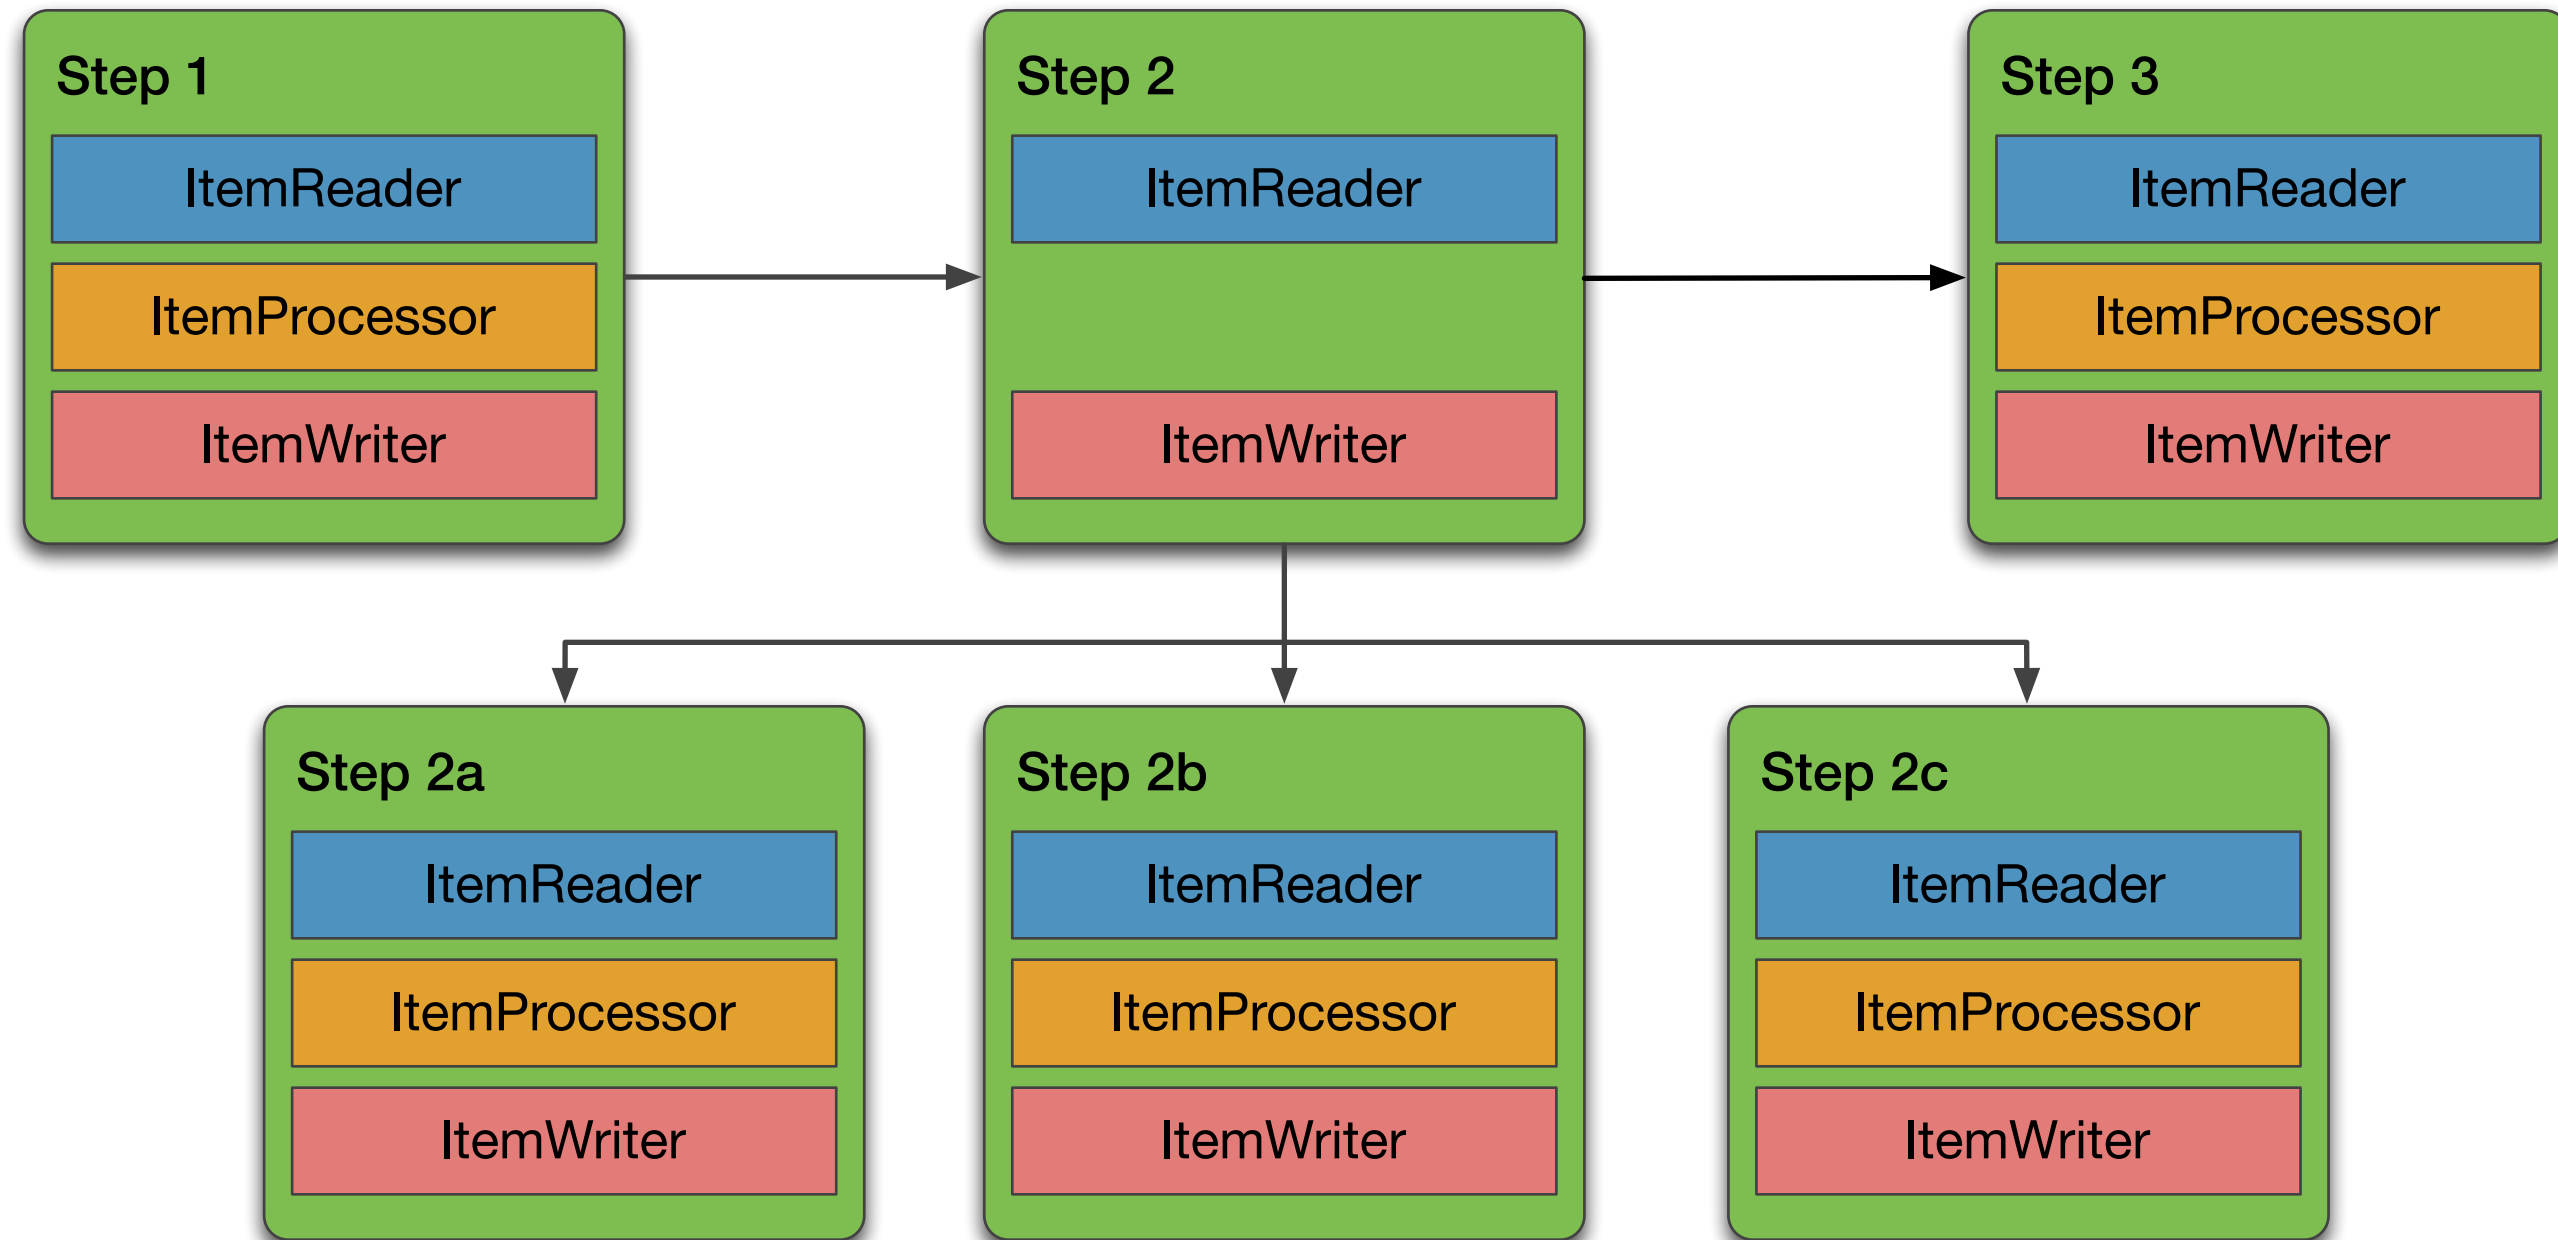
Multi-Threaded - Demo

Asynchronous Processors

- **AsyncItemProcessor**
 - Dispatches ItemProcessor logic on new thread, returning a Future to the AsyncItemWriter
- **AsyncItemWriter**
 - Writes the processed items after processing is complete

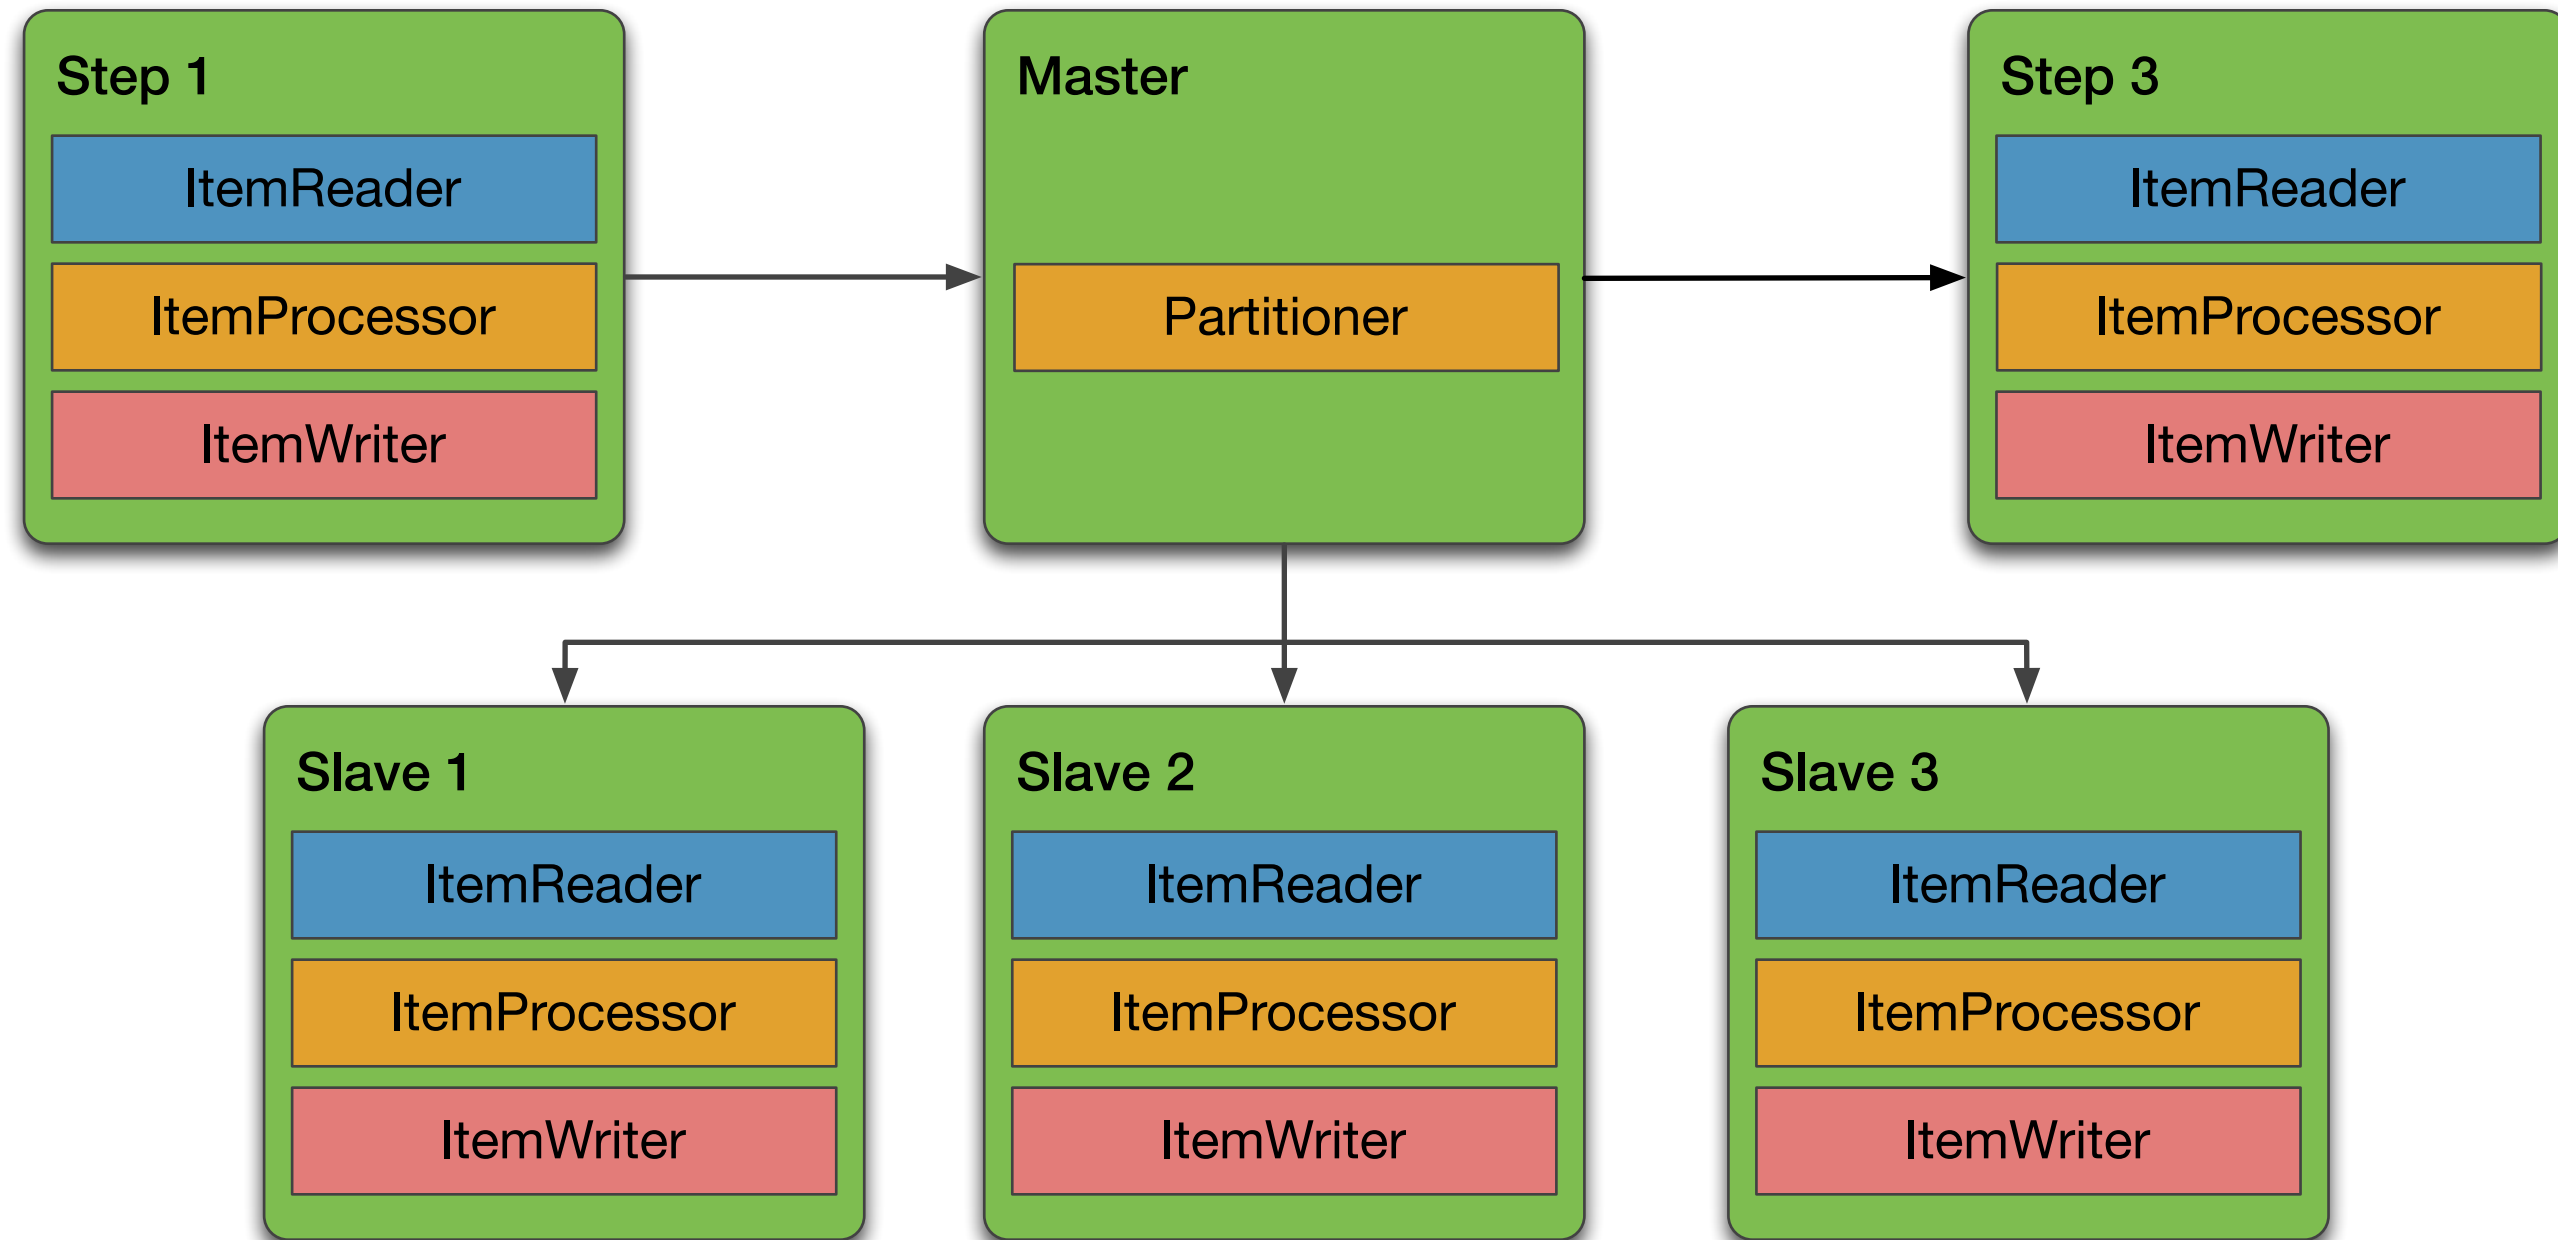
Asynchronous Processors - Demo

Remote Chunking



Remote Chunking - Demo

Remote Partitioning



Remote Partitioning - Demo

Demo - Launching via messages & informational messages

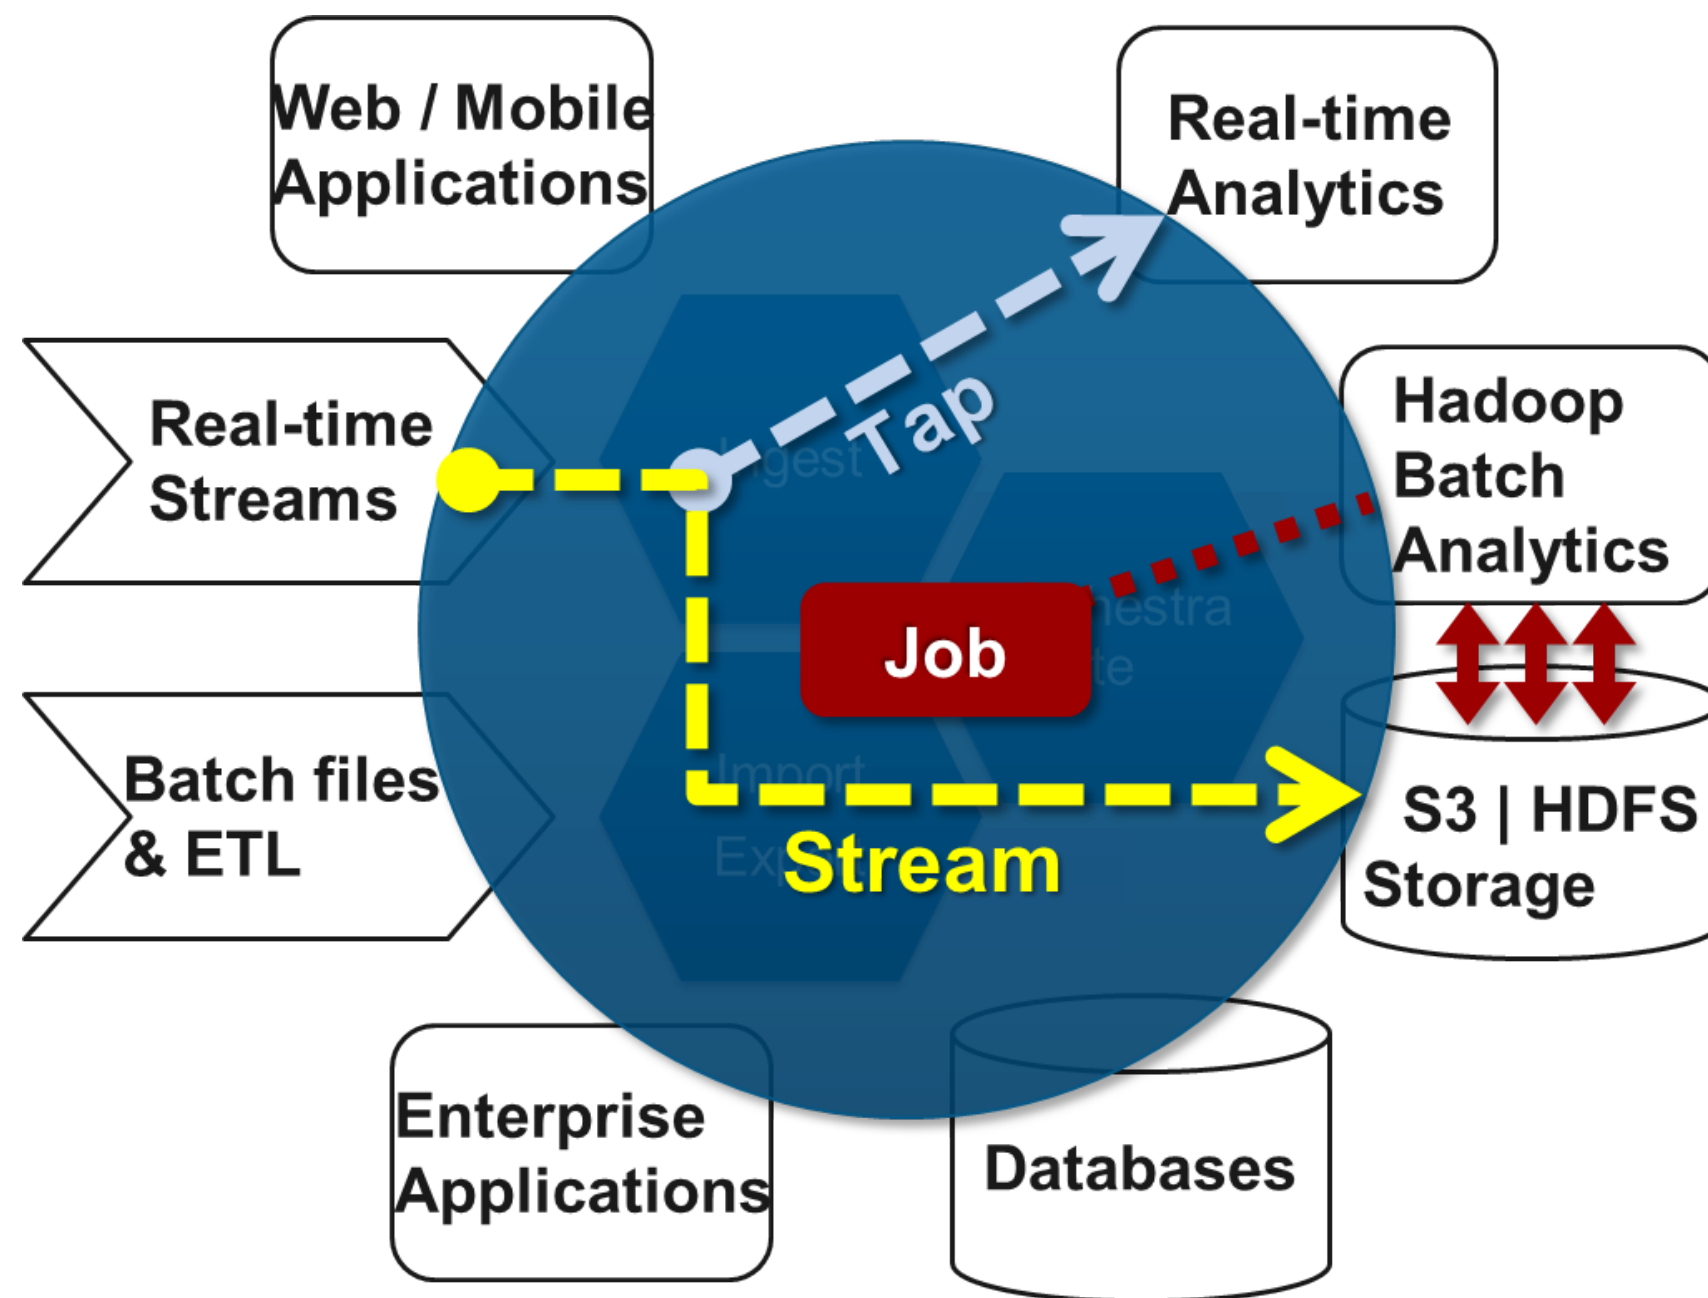
Does not provide scaling but demonstrates how launch job via messages and send information messages to integration points

Spring XD

<http://projects.spring.io/spring-xd/>

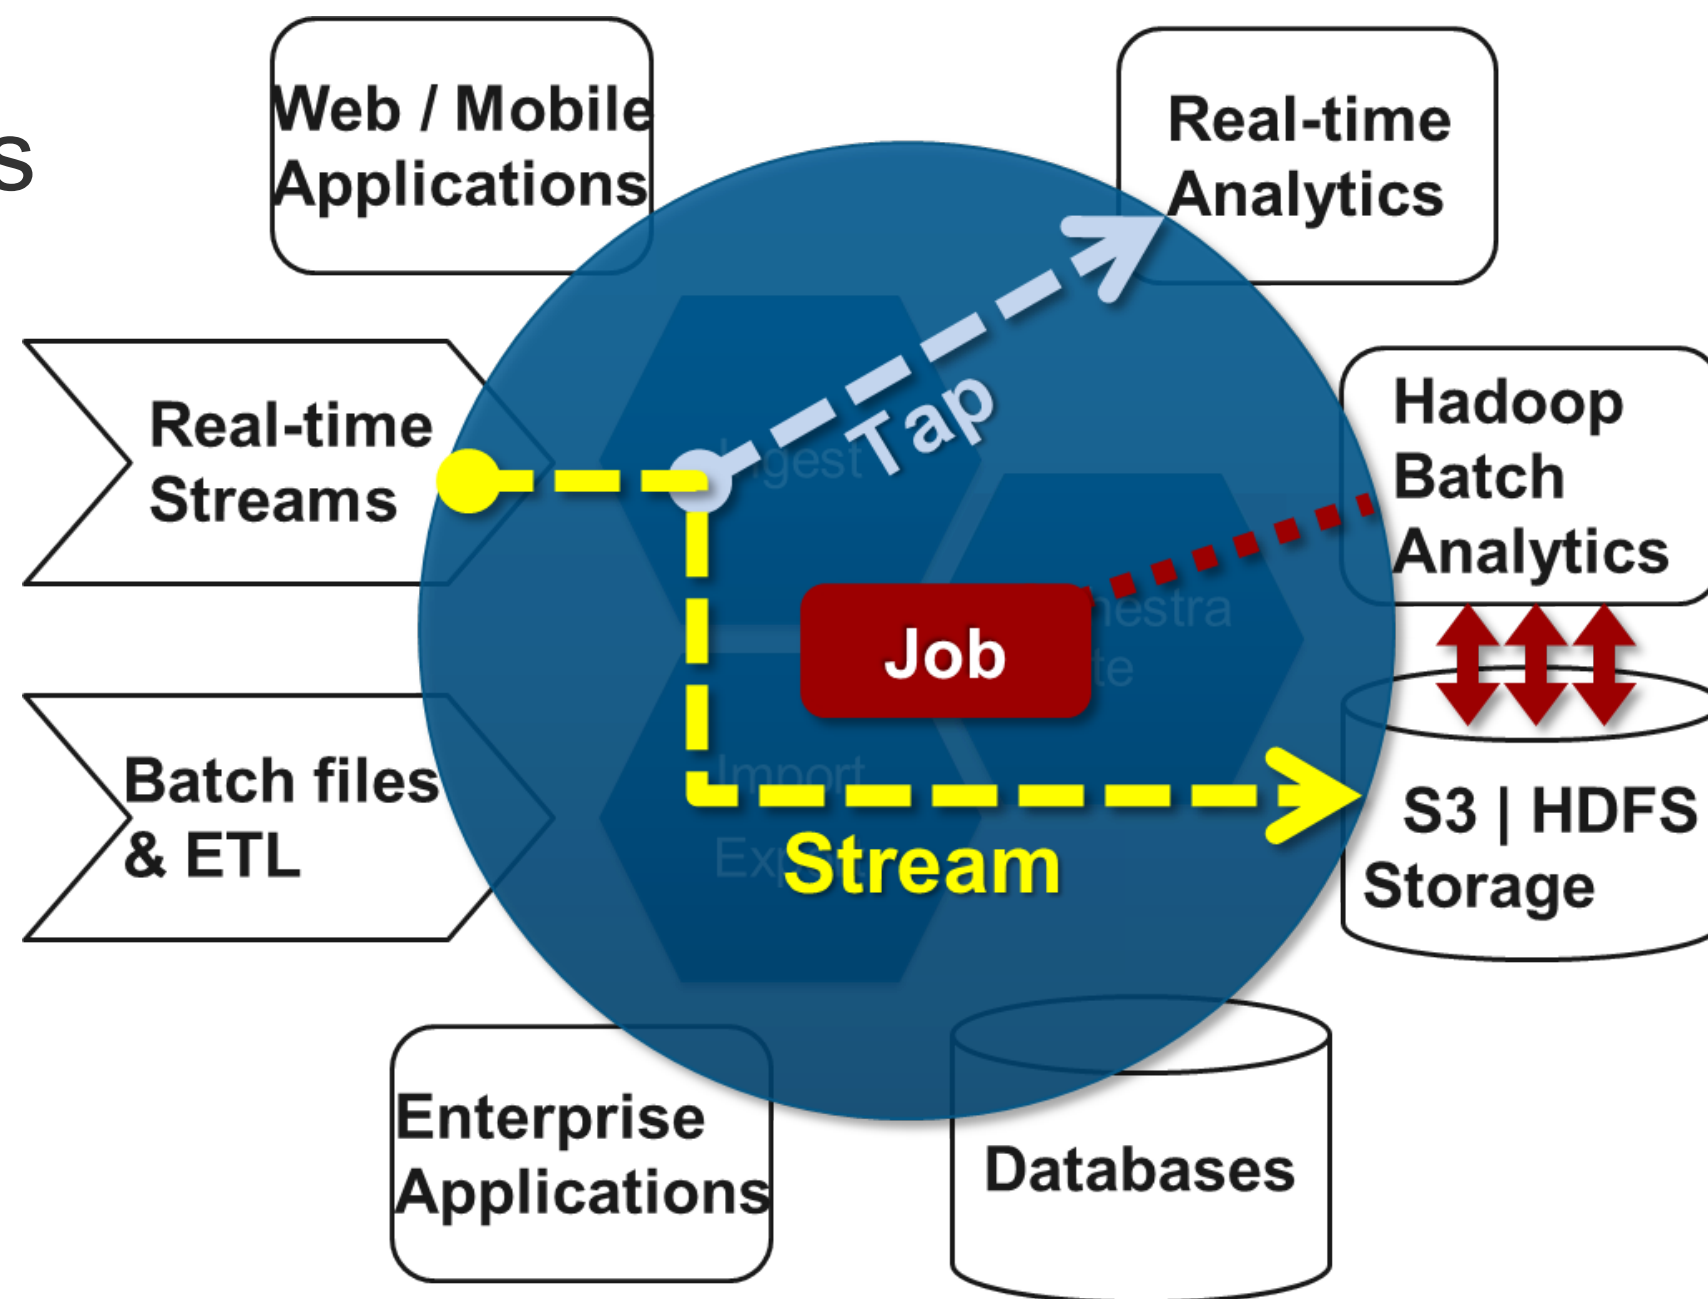
Tackling Big Data Complexity

- Data Ingestion
- Real-time Analytics
- Workflow Orchestration
- Data Export



Tackling Big Data Complexity cont.

- Built on existing Spring assets
 - Spring Integration
 - Spring Batch
 - Spring Data
 - Spring Boot
 - Spring for Apache Hadoop
 - Spring Shell
- Redis, GemFire, Hadoop



Data Ingestion Streams

- DSL based on Unix pipes and filters syntax

`http | file`

- Modules are parameterizable


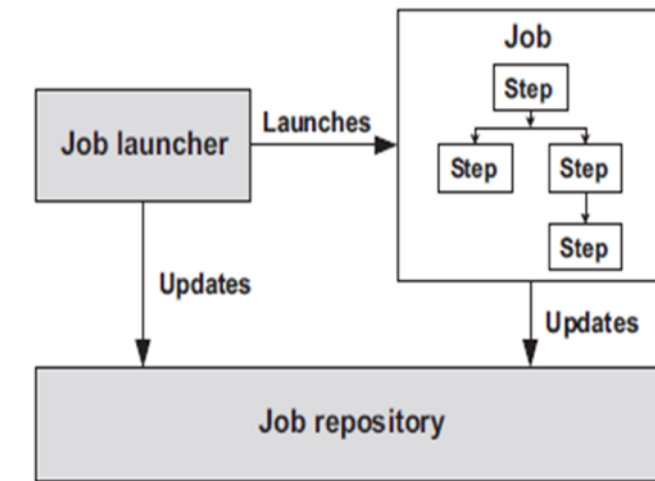
`twittersearch --query=spring | file --dir=/spring`

- Simple logic can be added via expressions or scripts

`http | filter --expression=payload=='Spring' | hdfs`


Hadoop workflow managed by Spring Batch

- Reuse Batch infrastructure and features to manage Hadoop workflows
- Job state management, launching, monitoring, restart/retry policies, etc.
- Step can be any Hadoop job type or HDFS script
- Can mix and match with other Batch readers/writers, e.g. JDBC for import/export use-cases



Manage Batch Jobs with Spring XD





STREAMS

JOBS

ABOUT

Create Definition for Job Module image-scaler

Definition Name

Demo

☒ Deploy created definition

Parameters

DateFormat

yyyy-MM-dd






Listeners

Job Execution Details - Execution ID: 4	
Property	Value
Id	4
Job Name	Demo
Job Instance	3
Job Parameters	random=0.3235123051449871
Start Time	2014-09-06 16:49:20,415
End Time	2014-09-06 16:49:23,043
Duration	2628 ms
Status	COMPLETED
	COMPLETED
	N/A
	2

ModulesDefinitionsDeploymentsExecutions			
Name	Execution Count	Last Execution Status	Actions
Demo	1	COMPLETED	<div><div>▶ Launch</div><div>⌚ Schedule</div><div>🔍</div></div>

Spring XD - Demo

Books



Learn More. Stay Connected.



THANK YOU!

Demo code and slides:

<https://github.com/SpringOne2GX-2014/spring-batch-performance-tuning>